

Causal Compression of Boolean Networks via Algorithmic Querying

Alberto Hernández-Espinosa

Hector Zenil

Abstract

We study the patterns that deterministic Boolean networks imprint across their exhaustive input-output repertoires. Gate rules, acting on a fixed connectivity structure, distribute information through the configuration space in a manner that is neither random nor immediately transparent: the resulting output sequences appear complex, yet conceal regularities that can be captured by short generative descriptions. We formalise this observation into a causal compression methodology situated in the exactly solvable limit of Algorithmic Information Dynamics—the regime in which the generating mechanism is given, so that the shortest description in a declared language may be computed rather than estimated—and grounded in the Principle of Computational Equivalence. For each of the twelve canonical Boolean gate families—monotone, parity, threshold, implication, negation, and canalising—we derive exact compressed representations and a deconvolution process that recovers the complete one-set of any gate without row-wise evaluation. We demonstrate compression first for isolated gates, then for multi-node queries where several gates are interrogated simultaneously, and finally for the full system, showing that the effective search space compresses by exact factors that depend on the interaction structure rather than the network size. Because the exact generating programme and its exact description length are both available here, these networks serve as a calibration class for methods that must operate without such knowledge: we measure a block-decomposition estimator against ground truth and find it within a factor of 2.49 of the true programme length—median 2.34 across 200 networks at the same size—while statistical baselines on the same object remain a factor of ≈ 43 away. These results demonstrate that the full complexity of a Boolean network’s behaviour is generated by short causal rules that can be identified, formalised, and exploited computationally.

1 Introduction

Boolean networks are deterministic dynamical systems, yet they exhibit the full phenomenology that makes complex systems singular: sensitive dependence on initial conditions, attractor multiplicity, and the exponential growth of the configuration space with system size. Since Kauffman’s pioneering work [5], they have served as objects of study in theoretical biology, gene regulation, and the broader investigation of how simple local rules generate complex collective behaviour [9].

Two fundamental questions pervade the analysis of such systems. First, what are the mechanisms by which order emerges in deterministic networks? Second, in what sense do global dynamics exhibit emergent properties that are not reducible to local update rules? These questions can be technically instantiated within Boolean networks, where each node applies a deterministic local rule and the full system behaviour is, in principle, computable. The difficulty is not computational power; it is *causal compression*. One can always enumerate a truth table or run a simulation. What remains elusive is a representation that is at once exact, compositional, and explanatory—one that identifies the generative rule structure responsible for the observed outputs rather than merely reproducing them [6, 9, 1, 7].

This paper presents a causal compression methodology for Boolean networks that is grounded in two intellectual traditions. The first is the Principle of Computational Equivalence (PCE) [9],

which holds that simple discrete rules generate the full complexity observed in natural and computational systems. The second is Algorithmic Information Dynamics (AID) [11, 13, 12], which provides an inferential and interventionist calculus for recovering short generative rules from observations, emphasising algorithmic probability, generative mechanisms, and causation over surface statistics.

Our relation to AID is best stated precisely, since it determines what may and may not be claimed here. AID in its general form addresses an inverse problem under uncertainty: the generating mechanism is unknown, and one estimates algorithmic content by perturbation and approximation of algorithmic probability. The present work occupies the corner of that programme in which the mechanism is *given*. A Boolean network is a controlled object—its wiring and its gates are known—so no estimation is required on the mechanism side, and the shortest description in a declared language can be written in closed form rather than approximated. We therefore present this method not as an alternative to AID but as its *exactly solvable limit*: the regime in which the quantities that AID must estimate can instead be computed. What we contribute back to AID is correspondingly specific—not a new estimator, but a class of objects with known exact generating programmes and known description lengths, against which estimators can be calibrated. We exercise that calibration in Section 5.2.

The currency of the argument follows from this position. A description is taken to be explanatory when it is a programme that reproduces the observed behaviour exactly, and its cost is measured by the length of that programme rather than by the entropy of the behaviour it produces. This is the sense in which compression stands in for comprehension [10]: a short programme that regenerates a large object constitutes an account of that object, whereas an entropy figure quantifies a coding cost and explains nothing. The distinction is not rhetorical here but measurable, and Section 5.2 reports the measurement.

Our analysis is based on the observation that when the output of a Boolean network is examined node by node across the ordered exhaustive repertoire, the isolated output sequences are not random: they exhibit self-similar patterns replicated at several scales. The resulting information distribution—what was termed “fractal behaviour” in earlier work [3]—is a structural consequence of the interaction between the network topology and the deterministic local rules.

The key insight is that this self-similar structure admits exact compressed representations. For any Boolean gate operating within a network of n nodes, the complete set of repertoire indices at which the gate outputs one can be expressed as two short objects: a *base set* L of decimal anchors, and an *offset family* Ω . Deconvolution of this pair—adding each sumando to each decimal anchor—recovers the exact one-set without evaluating a single row. The base set encodes the gate semantics; the offset family encodes the network topology. Together, they constitute a short generative programme whose output is the full behaviour of the gate across 2^n input states.

The self-similarity is not a metaphor but a property of Ω , and it can be stated mechanically. Each coordinate that does not feed the node contributes an independent binary choice, so Ω factorises as a sumset of powers of two, $\bigoplus_j \{0, 2^{j-1}\}$ over the free coordinates. The activation pattern therefore recurs at every partial sum of those weights, which is dyadic rescaling in a literal sense: the same short pattern, laid down at geometrically spaced intervals across the repertoire. This is what was observed informally as “fractal” behaviour, and it is exactly why a description of the pair (L, Ω) suffices where a listing of the sequence does not.

In this paper we treat twelve families: AND, OR, XOR, NAND, NOR, XNOR, NOT, IMPLIES, NIMPLIES, KOFN (threshold), MAJORITY, and CANALISING, which exhaust the standard gate catalogue used in Boolean network modelling.

We compare experiments with networks of different sizes to show that the compression factor is independent of the network size and we confirm that analytic evaluation time does not grow with the ambient dimension n —it is governed by the support-union instead, and stays of the order of a millisecond or below across the whole range tested.

The perspective is *causality-first* in three senses. First, it is rule-first rather than correlation-first: the object of study is the generative logic of the system, not its observed frequencies. Second, it is deterministic rather than statistical: the method reconstructs exact behaviour under known dynamics, not approximate inference under unknown mechanisms. Third, it is compositional: global behaviour is built from local causal rules, and the compression reflects the causal structure of the mechanism itself. All computational experiments are reproducible via the companion code repository (see Appendix A).

2 Background

2.1 Boolean networks and ordered repertoires

A Boolean network of size n is specified by a connectivity matrix $cm \in \{0, 1\}^{n \times n}$, where $cm_{ij} = 1$ means node j feeds node i , and a dynamic vector $dyn = (d_1, \dots, d_n)$ listing the local Boolean gate at each node (see lines 1-11 in Listing code 2).

The connected-input set of node i is $N_i = \{j : cm_{ij} = 1\}$, and the synchronous one-step update is

$$F(x) = (g_1(x_{N_1}), \dots, g_n(x_{N_n})), \quad x \in \{0, 1\}^n.$$

where x represents an input vector in the input repertoire that composes an *ordered exhaustive repertoire* that is the set of all 2^n input vectors under a declared enumeration. Throughout this paper we use LSB-first ordering, in which the j -th vector is $v(j) = \text{Reverse}(\text{IntegerDigits}(j-1, 2, n))$ with bit weights $w(i) = 2^{i-1}$. The resulting index universe is $\mathcal{U}_n = \{1, \dots, 2^n\}$.

2.2 Gate families and output signatures

When the output of a single node is isolated across the ordered exhaustive repertoire, the resulting binary sequence, this is, a string of ones and zeros of length 2^n , reveals patterns characteristic of the gate that node implements. These patterns are the starting point of the method: before any formal derivation, one observes the regularities in the output sequences that each gate family produces when embedded in a network.

The observation is most transparent for a small case examined in full. Consider the seven-node network of Listing 1, and within it node 4: an AND gate reading inputs $\{1, 3, 5, 7\}$.

Listing 1: Seven-node network used throughout this section. Node 4 is an AND gate over inputs $\{1, 3, 5, 7\}$; nodes $\{2, 4, 6\}$ do not feed it.

```

1 cm07 = {{0, 0, 1, 0, 0, 0, 1},
2         {0, 0, 1, 0, 0, 1, 0},
3         {1, 0, 0, 0, 1, 0, 1},
4         {1, 0, 1, 0, 1, 0, 1},
5         {0, 0, 1, 1, 0, 1, 1},
6         {1, 1, 1, 0, 0, 0, 0},
7         {0, 1, 0, 1, 1, 1, 0}};
8 dyn07 = {"AND", "OR", "OR", "AND", "OR", "OR", "AND"};
```

Isolating node 4 across the ordered exhaustive repertoire gives a 128-element binary string. It is dominated by zeros, since an AND gate fires only when *all* connected inputs are simultaneously active. Written out in full, with the eight firing positions marked, the string is

```

positions 0 ..... 84 (all 0)
positions 85 86 87 88 ... 92 93 94 95 96 ..... 116
           1  0  1  0       1  0  1  0              (0)
positions 117 118 119 120 ... 124 125 126 127
           1  0  1  0       1  0  1
```

so the gate outputs one exactly at positions $\{85, 87, 93, 95, 117, 119, 125, 127\}$, using 0-based repertoire indices.

A flat run-length encoding of this string reads $\{85 \rightarrow 0, 1 \rightarrow 1, 1 \rightarrow 0, 1 \rightarrow 1, 5 \rightarrow 0, 1 \rightarrow 1, 1 \rightarrow 0, 1 \rightarrow 1, 21 \rightarrow 0, \dots\}$, where each pair $k \rightarrow v$ denotes k consecutive occurrences of digit v . Sixteen such runs, thirty-two tokens, describe the sequence completely.

That encoding is nevertheless the wrong object, and it is worth saying why plainly, because the distinction is the whole point of the method. A run-length listing is a *tally*: it records the sequence faithfully and generates nothing. It offers no account of why the gaps between firings have lengths 1, 5 and 21 rather than any other lengths, and it would grow without bound as n grows. What the method seeks is the rule that produces those gaps.

That rule is exact and short. The four connected inputs carry bit weights $1+4+16+64 = 85$, and an AND gate fires only when all four are set, which fixes a single decimal anchor: the base set is $L = \{85\}$. The three disconnected coordinates $\{2, 4, 6\}$ carry weights $\{2, 8, 32\}$ and are free to take any values, so the offset family is the set of all their subset sums. Crucially, that family *factorises*, one independent binary choice per free coordinate:

$$\Omega = \{0, 2\} \oplus \{0, 8\} \oplus \{0, 32\} = \{0, 2, 8, 10, 32, 34, 40, 42\},$$

where $\oplus$ denotes the sumset $A \oplus B = \{a+b\}$. Deconvolution recovers the firing positions exactly:

$$\text{Dec}(L, \Omega) = \{85, 87, 93, 95, 117, 119, 125, 127\}.$$

Two features of this description deserve emphasis. First, it is *nested* rather than sequential: the activation pattern is stated once and recurs at every partial sum of the free weights, which is exactly why the sequence appears self-similar at dyadic scales, and exactly what the flat tally destroys. The gap lengths 1, 5 and 21 are not primitive data but consequences of the weights 2, 8 and 32. Second, it is dramatically shorter: one decimal anchor and three free weights—four numbers—against thirty-two run-length tokens, and the gap widens as n grows because L and Ω depend on the node’s support rather than on the size of the repertoire. Every figure in this example is regenerated by the companion script `worked_example_7node.py`.

An OR gate produces the complementary picture: long runs of ones interrupted by rare clusters of zeros, since the gate fires whenever *any* input is active. XOR and XNOR gates yield periodic alternation governed by the parity of active inputs. Threshold gates (KOFN, MAJORITY) produce layered transitions whose boundaries correspond to Hamming-weight levels. Implication and canalising gates exhibit asymmetric signatures reflecting the privileged role of specific inputs.

The representation is a schema normal form. The structure just exhibited has a precise name. A *schema* in Holland’s sense [4] is a template over $\{0, 1, *\}$, in which $*$ marks a position whose value does not affect membership—and it is worth noting that Holland’s $*$ is defined per template, which is precisely why *free* must be read per schema here as well. The worked example is exactly such a template: node 4 fires on the schema $1 * 1 * 1 * 1$ read over coordinates 1 to 7, with its four connected inputs $\{1, 3, 5, 7\}$ fixed at one and the three disconnected coordinates $\{2, 4, 6\}$ free. This is not an analogy. For every gate family in Table 1, the one-set is precisely a union of schemata, and the pair (L, Ω) is a normal form for that union: L enumerates the determined coordinates, Ω enumerates the fillings of the don’t-care positions. Deconvolution is the instantiation of the template.

The correspondence is quantitative as well as structural. The *order* of a schema—its number of fixed positions—is exactly the support size $|C_q|$ of the corresponding query, and the count of don’t-care positions is $n - |C_q|$, whence $|\Omega| = 2^{n-|C_q|}$. The scalability result reported in Section 5.1, that analytic cost is governed by $|C_q|$ rather than by n , is therefore a statement about schema order: low-order schemata are cheap to evaluate however large the ambient space becomes.

What the disconnected coordinates do. The disconnected coordinates are the subset of the don't-cares that is free in *every* schema of a node, and it is worth being exact about their role, because the natural summary—that they are irrelevant—is true at only one of three levels. Nothing in what follows should be read as a definition of Ω ; it is the behaviour of one always-present part of it.

- *Per node, per step.* They do not affect the output *value*. This is precisely what makes them don't-care positions, and it is why the gate semantics can be read off from L alone.
- *Per repertoire.* They nevertheless determine *where* and *how often* the activation pattern occurs: their bit weights generate Ω , fixing its cardinality $2^{n-|C_q|}$ and the dyadic spacing at which the pattern recurs. They are causally inert as to the value and constitutive of the structure of the behaviour.
- *Per query.* Membership of the don't-care set is not intrinsic to a node. When several nodes are interrogated jointly, the free set is recomputed against the *union* of their supports, so a coordinate that is a don't-care for one node ceases to be one as soon as a co-queried node reads it. The amount by which the free set contracts is the overlap μ_q of Section 4.3, and the resulting factor 2^{μ_q} is what makes joint queries cheaper than the product of separate ones.

The third level is where interaction enters. Don't-care status is a property of the query, not of the network, and integration is measured by exactly how much it shrinks under joint interrogation. We note that this concerns the geometry of the one-set under one synchronous step; over several steps a coordinate outside C_q may of course become an ancestor of the node, which is a separate question.

Two observational tools proved essential for identifying these regularities. First, *run-length encoding* compresses each output sequence into a short list of $(k \rightarrow v)$ pairs, making the repetitive structure visible—though, as the worked example showed, it records rather than generates. Second, schemata identify which input coordinates determine the output and which are free to vary. Together, these tools revealed that every gate family produces output sequences whose structure can be described by a short generative rule, and that the formal characterisation of each family follows from formalising these observed patterns.

Table 1 summarises the formal one-set structure—the set of repertoire indices at which each gate outputs one—that emerges from this pattern analysis, expressed in terms of coordinate bands, parity classes, and Hamming-weight layers.

Gate family	One-set structure
AND	Intersection of one-bands of all connected inputs
OR	Union of one-bands
NAND	Complement of AND one-set
NOR	Intersection of zero-bands (complement of OR)
XOR	Odd-parity class over connected inputs
XNOR	Even-parity class
NOT	Zero-band of the single connected input
IMPLIES	Complement of (antecedent one-band $\cap$ consequent zero-band)
NIMPLIES	Antecedent one-band $\cap$ consequent zero-band
KOFN	Union of Hamming-weight layers at or above threshold k
MAJORITY	KOFN with $k = \lceil d/2 \rceil$
CANALISING	Trigger set $\cup$ (non-trigger $\cap$ fallback one-set)

Table 1: Formal one-set structure for each gate family, derived from the observed output patterns. Band $\mathcal{B}_i^a$ denotes the set of indices where coordinate i equals a ; parity and weight classes are defined analogously. Full derivations appear in [2].

The key observation is that *all* of these one-sets share a common structural property: each decomposes into a short activation pattern that encodes the gate semantics, replicated across the repertoire by shifts that encode the network topology. This universal decomposition is what makes compressed representation possible across the full gate catalogue.

2.3 Self-similar patterns and compressed representation

The patterns described above are not merely compact descriptions of individual outputs; they reveal a deeper structural regularity. When the output of a single node is examined across the ordered exhaustive repertoire, the same activation sub-pattern recurs at geometrically spaced intervals throughout the sequence. This self-similar structure—what was termed “fractal behaviour” in earlier work [3]—is not an artefact of any particular gate type but a general property of deterministic Boolean networks operating on binary-encoded input spaces.

The discovery followed a natural path. The perturbation test—evaluating the network on every possible input—produces for each node an output sequence of length 2^n . Examined in isolation, these sequences appear complex: for a seven-node network, the output of any given node is a 128-element binary string with no immediately obvious rule. Yet when described using run-length notation and schemata, the same short sub-pattern is found to replicate across the sequence at intervals that follow a geometric progression tied to powers of two.

This observation admits a precise representation as two short objects:

- The **base set** L (termed *DecimalRepertoire* in the companion code): the repertoire indices at which the node outputs one when only its connected inputs vary and all other coordinates are held at zero. This set captures the gate’s intrinsic activation rule as it operates within the network.
- The **offset family** Ω (termed *Sumandos*): the set of shifts at which the base-set pattern is replicated across the full repertoire. Each shift is a filling of the *don’t-care positions of the schema in question*, and the geometric spacing of these shifts is what produces the self-similar structure observed in the output sequences.

A point of definition that must be stated precisely, because the imprecise form is easy to adopt and hard to unlearn. A coordinate is *free* relative to a schema, not relative to a node. The disconnected coordinates are free in every schema and are therefore always present in Ω , but they do not exhaust it: a connected input may be free in some schemata and fixed in others. Rule 110 makes the distinction concrete. All three of its inputs are connected, so on the narrow reading its free set is empty, $\Omega = \{0\}$, and L must list all five minterms; on the correct reading it is three schemata, $01*$, $10*$ and $*10$, every one of whose don’t-cares sits on a *connected* input. The distinction is quantitative and not merely verbal: it is what separates an OR of in-degree ten, whose 1023 minterms merge into ten schemata, from an XOR of the same in-degree, whose 512 minterms merge into none.

The full one-set is recovered by *deconvolution*:

$$\mathcal{J} = \text{Dec}(L, \Omega) = \{\ell + \omega : \ell \in L, \omega \in \Omega\}.$$

This operation unfolds the compressed pair into the exact set of repertoire positions at which the node outputs one, without evaluating a single row of the truth table. The base set encodes the short rule that the gate produces; the offset family encodes where that rule repeats.

The mathematical explanation for this decomposition is straightforward: the connected inputs determine the gate’s activation condition, while the disconnected inputs—those coordinates absent from the node’s row in the connectivity matrix—are free to take any value without affecting the output. Their bit weights generate the always-present part of the offset family as a sumset of powers of two; the remainder of Ω comes from connected inputs that the individual

schema does not constrain. But the method was not derived from this algebraic observation. It was discovered by inspecting the output sequences, identifying their self-similar structure, and then recognising that the structure could be expressed as a short generative rule. The connected/disconnected decomposition is the mathematical explanation of *why* the patterns exist; the patterns themselves are what the method observes and exploits.

2.4 Causal querying and algorithmic complexity

Within the AID framework [11, 13], the compressed representation constitutes a short generative programme for the gate’s behaviour, and the deconvolution is its execution. The methodology is a form of causal querying: given a known deterministic mechanism (the Boolean network), one asks which input conditions produce a specified output event, and the answer takes the form of a programme that is dramatically shorter than the output it generates.

The object described is accordingly individual rather than ensemble [3]. Methods built on regression, correlation or probability distributions—transfer entropy, Granger causality, partial information decomposition—characterise the family of sequences a source emits. The base set and offset family instead describe this network’s input–output map and no other. Section 5.2 makes the consequence quantitative.

This connects to the perturbation asymmetry central to AID [11]. An algorithmically simple object—one generated by a short programme—is sensitive to perturbation: modifying the programme propagates non-linearly through the output. A random object is insensitive: its shortest description is the object itself. The compressed representation of a Boolean network is exactly such a short programme, and its brevity relative to the full output table is a quantitative signature of the causal structure of the mechanism.

3 Causal Querying and Compressed Representations

All code listings in this section and the next use two core functions from the companion library (see Appendix A): `ApplyGate[gate, inputs, params]` evaluates a single Boolean gate on a given input vector, and `CreateRepertoiresDispatch[cm, dyn, params]` builds the full 2^n -row exhaustive repertoire by applying `ApplyGate` to every row of the ordered input space.

3.1 Single-gate case: AND deconvolution

We begin with the simplest non-trivial case: a single AND gate inside a six-node network.

Listing 2: Six-node network definition and exhaustive repertoire via gate dispatch.

```

1  (* Load companion library *)
2
3  cm06 = {
4      {1, 0, 0, 0, 0, 0},
5      {0, 1, 0, 0, 0, 0},
6      {0, 0, 1, 0, 0, 0},
7      {1, 0, 0, 1, 0, 0},
8      {0, 1, 0, 1, 0, 0},
9      {1, 0, 1, 0, 1, 0}
10 };
11 dyn06 = {"OR", "NOT", "OR", "IMPLIES", "AND", "XOR"};
12
13 dispatch06 = CreateRepertoiresDispatch[cm06, dyn06];
14 inputs06 = dispatch06["RepertoireInputs"];
15 outputs06 = dispatch06["RepertoireOutputs"];
16
17 Length[inputs06] = 64

```

```

18 inputs06[[1;;3]] = {{0,0,0,0,0,0},{1,0,0,0,0,0},{0,1,0,0,0,0}}
19 outputs06[[1;;3]] = {{0,1,0,1,0,0},{1,1,0,0,0,1},{0,0,0,1,0,0}}

```

Node 5 is an AND gate with connected-input set $N_5 = \{2, 4\}$, extracted directly from the connectivity matrix. The network has $n = 6$, so the exhaustive repertoire contains $2^6 = 64$ rows. The question is: at which of these 64 rows does node 5 output one?

The compressed representation answers this without evaluating any row. Under LSB-first ordering, the bit weights are $w(i) = 2^{i-1}$, so the connected contribution (the *decimal anchor*) is $P(N_5) = w(2) + w(4) = 2 + 8 = 10$. The disconnected inputs are $\{1, 3, 5, 6\}$, and their offset family is the set of all sums of subsets of the corresponding weights:

Listing 3: Computing the offset family and AND deconvolution.

```

1 weights[n_] := 2^Range[0, n- 1];
2
3 allOffsets[n_, connected_] := Module[
4   {free = Complement[Range[n], connected], ws},
5   ws = weights[n][[free]];
6   Sort[(# . ws) & /@ Tuples[{0, 1}, Length[ws]]]
7 ];
8
9 givePlaces[locations_, sumandos_] :=
10   Sort@Flatten[Table[loc +sumandos, {loc, locations}]];
11
12 ic5 = Sort@Flatten@Position[cm06[[5]], 1]
13 decimalAnchor5= 1+ Total[weights[6][[ic5]]]
14 offsets5= allOffsets[6, ic5]
15 predicted5= givePlaces[{decimalAnchor5}, offsets5]

```

```

(* ic5      = {2, 4}
(* decimalAnchor5 = 11
(* offsets5 = {0,1,4,5,16,17,20,21,32,33,36,37,48,49,52,53} *)
(* predicted5 = {11,12,15,16,27,28,31,32,43,44,47,48,59,60,63,64} *)

```

Corroboration against the exhaustive baseline produced by `CreateRepertoiresDispatch`:

Listing 4: Verification against exhaustive baseline.

```

1 baseline5= Flatten@Position[outputs06[[All, 5]], 1]
2 Sort[predicted5] === Sort[baseline5]

```

```

(* baseline5 = {11,12,15,16,27,28,31,32,43,44,47,48,59,60,63,64} *)
(* Out: True

```

The analytic formula and the exhaustive gate-dispatch table agree on all 64 rows. The 16-element one-set is fully determined by two short objects: one decimal anchor $\{11\}$ and one 16-element offset family. The rule itself is encoded by a single decimal anchor; the offsets are a structural consequence of the network topology, shared across all gates with the same disconnected coordinates.

Table 2 shows the correspondence between free-bit assignments, offsets, and repertoire indices.

3.2 Extension to all gate families

The `indexSetAnalytic` function computes the analytic one-set for any of the twelve gate families. The dispatch is purely structural: each gate family determines a selection rule over the connected-input assignments, and the free coordinates generate the offset family uniformly.

Free bits (v_1, v_3, v_5, v_6)	Offset δ	Index $j = 11 + \delta$	AND output
(0, 0, 0, 0)	0	11	1
(1, 0, 0, 0)	1	12	1
(0, 1, 0, 0)	4	15	1
(1, 1, 0, 0)	5	16	1
(0, 0, 1, 0)	16	27	1
(1, 0, 1, 0)	17	28	1
(0, 1, 1, 0)	20	31	1
(1, 1, 1, 0)	21	32	1
(0, 0, 0, 1)	32	43	1
(1, 0, 0, 1)	33	44	1
(0, 1, 0, 1)	36	47	1
(1, 1, 0, 1)	37	48	1
(0, 0, 1, 1)	48	59	1
(1, 0, 1, 1)	49	60	1
(0, 1, 1, 1)	52	63	1
(1, 1, 1, 1)	53	64	1

Table 2: AND deconvolution on the six-node network. The decimal anchor $P(N_5) = 10$ gives base index 11. Each row corresponds to one assignment of the four free bits, producing an offset δ and an index $j = 11 + \delta$ at which the AND gate is active.

Listing 5: Complete dispatch of `indexSetAnalytic` for all gate families.

```

1 indexSetAnalytic[n_Integer, Ic_List, gate_String,
2   params_Association : <|>] := Module[
3   {free, pow, d, indexFromPos, subsFree,
4   k, strict, pair, a, b, ii},
5   free = Complement[Range[n], Ic];
6   pow = weights[n]; d = Length[Ic];
7   indexFromPos[pos_List] := 1 + Total[pow[[pos]]];
8   subsFree = Subsets[free];
9   Which[
10    gate === "AND",
11      indexFromPos[Join[Ic, #]] & /@ subsFree,
12    gate === "OR",
13      Complement[Range[1, 2^n],
14        indexFromPos/@ subsFree],
15    gate === "XOR",
16      Module[{assigns},
17        assigns = Select[Tuples[{0, 1}, d],
18          Mod[Total[#], 2] == 1 &];
19        Flatten@Table[indexFromPos[
20          Join[Pick[Ic, assigns[[u]], 1], s]],
21          {u, Length[assigns]}, {s, subsFree}]],
22    gate === "XNOR",
23      Module[{assigns},
24        assigns = Select[Tuples[{0, 1}, d],
25          Mod[Total[#], 2] == 0 &];
26        Flatten@Table[indexFromPos[
27          Join[Pick[Ic, assigns[[u]], 1], s]],
28          {u, Length[assigns]}, {s, subsFree}]],
29    gate === "NAND",
30      Complement[Range[1, 2^n],
31        indexFromPos[Join[Ic, #]] & /@ subsFree],
32    gate === "NOR",
33      indexFromPos/@ subsFree,

```

```

34 gate === "NOT",
35   ii = Lookup[params, "i", First[Ic]];
36   indexFromPos/@
37     Subsets[Complement[Range[n], {ii}]],
38 gate === "IMPLIES",
39   pair = Lookup[params, "pair", Ic[;; 2]];
40   a = pair[[1]]; b = pair[[2]];
41   Complement[Range[1, 2^n],
42     indexFromPos/@ (Join[{a}, #] & /@
43       Subsets[Complement[Range[n], {a, b}]]),
44 gate === "NIMPLIES",
45   pair = Lookup[params, "pair", Ic[;; 2]];
46   a = pair[[1]]; b = pair[[2]];
47   indexFromPos/@ (Join[{a}, #] & /@
48     Subsets[Complement[Range[n], {a, b}]]),
49 gate === "MAJORITY",
50   Module[{t, assigns},
51     t = Floor[d/2] + 1;
52     assigns = Select[Tuples[{0, 1}, d],
53       Total[#] >= t &];
54     Flatten@Table[indexFromPos[
55       Join[Pick[Ic, assigns[[u]], 1], s]],
56       {u, Length[assigns]}, {s, subsFree}],
57 gate === "KOFN",
58   Module[{assigns},
59     k = Lookup[params, "k", 1];
60     strict = TrueQ[Lookup[params,
61       "strict", False]];
62     assigns = Select[Tuples[{0, 1}, d],
63       If[strict, Total[#] > k,
64         Total[#] >= k] &];
65     Flatten@Table[indexFromPos[
66       Join[Pick[Ic, assigns[[u]], 1], s]],
67       {u, Length[assigns]}, {s, subsFree}],
68   True, {}
69 ]
70 ];

```

As a verification, the AND one-set from `indexSetAnalytic` matches the deconvolution above exactly:

```
1 Sort@indexSetAnalytic[6, {2, 4}, "AND"]
```

```
(* {11,12,15,16,27,28,31,32,43,44,47,48,59,60,63,64} *)
```

The structural pattern across gate families is as follows:

- **Monotone gates** (AND, OR, NAND, NOR): the base set is determined by band intersections or their complements. AND has a single decimal anchor; OR has $2^d - 1$ decimal anchors.
- **Parity gates** (XOR, XNOR): the base set contains all connected-input assignments satisfying the parity condition. The number of decimal anchors is exactly 2^{d-1} .
- **Threshold gates** (KOFN, MAJORITY): the base set is the union of Hamming-weight layers at or above the threshold.
- **Implication gates** (IMPLIES, NIMPLIES): asymmetric; the base set depends on a designated antecedent–consequent pair.

- **Negation** (NOT): the base set is the zero-band of the single input, with 2^{n-1} elements.
- **Canalising**: hierarchical collapse—the trigger condition determines a primary set, and the fallback rule handles the remainder.

In all cases, the deconvolution mechanism is uniform: the gate-specific base set is combined with the topology-determined offset family to produce the exact one-set. All twelve families are verified computationally in the 10-node benchmark that follows.

3.3 Ordering conventions

The bit-reversal involution $\varphi(j) = 1 + \text{FromDigits}(\text{Reverse}(\text{IntegerDigits}(j-1, 2, n)), 2)$ transports all results between LSB-first and MSB-first orderings. Computational verification on the six-node benchmark confirms that every analytic one-set is preserved under φ : if $\mathcal{J}$ is the one-set under one ordering, then $\varphi(\mathcal{J})$ is the one-set under the other. The formal proof of this invariance appears in [2]; here we note that no numerical anchor in this paper depends on the ordering choice, since φ is a bijection on $\mathcal{U}_n$.

4 Mixed-Gate Networks and Overlap Compression

4.1 The 10-node benchmark

To demonstrate the method on a heterogeneous network, we define a 10-node benchmark containing one instance of each non-trivial gate family. The exhaustive repertoire is built by `CreateRepertoiresDispatch`, which calls `ApplyGate` on every row:

Listing 6: Ten-node mixed-gate network definition and repertoire.

```

1 cm10 = {
2   {0,1,1,0,0,0,0,0,0,0},
3   {1,0,1,0,0,0,0,0,0,0},
4   {0,0,0,1,1,0,0,0,0,0},
5   {0,1,1,0,1,0,0,0,0,0},
6   {0,0,0,0,0,1,0,0,0,0},
7   {0,0,0,0,1,0,1,0,0,0},
8   {0,0,0,0,0,1,0,0,0,0},
9   {1,0,0,0,0,0,0,0,1,0},
10  {0,1,0,0,0,0,0,0,0,1},
11  {0,0,1,1,0,0,1,1,0,0}
12 };
13 dyn10 = {"AND", "OR", "XOR", "KOFN", "NOR",
14          "XNOR", "NOT", "IMPLIES", "NIMPLIES", "MAJORITY"};
15 params10 = <|4 -> <|"k" -> 2|>,
16            8 -> <|"pair" -> {1, 9}|>,
17            9 -> <|"pair" -> {2, 10}|>|>;
18
19 dispatch10 = CreateRepertoiresDispatch[cm10, dyn10, params10];
20 inputs10 = dispatch10["RepertoireInputs"];
21 outputs10 = dispatch10["RepertoireOutputs"];

```

```
(* Length[outputs10] = 1024 *)
```

The network is depicted in Figure 1. Node 4 is a 2-of-3 threshold gate (KOFN with $k = 2$), node 10 is a strict majority gate on four inputs, and nodes 8 and 9 have designated antecedent–consequent pairs for their implication gates.




Figure 1: Ten-node mixed-gate benchmark network. Blue nodes use monotone or parity gates; red nodes use threshold, negation, implication, or majority gates.

The connected-input sets are extracted from the connectivity matrix, and each node’s analytic one-set is computed via `indexSetAnalytic` and verified against the exhaustive baseline on all $2^{10} = 1024$ rows:

Listing 7: Per-node verification against exhaustive baseline.

```

1  ics10= Table[Sort@Flatten@Position[cm10[[k]], 1],
2      {k, 1, 10}];
3  oneSets10= Table[
4      Sort@indexSetAnalytic[10, ics10[[k]],
5      dyn10[[k]], Lookup[params10, k, <||>]],
6      {k, 1, 10}];
7  baselineOneSets10= Table[
8      Flatten@Position[outputs10[[All, k]], 1],
9      {k, 1, 10}];
10
11 And @@ Table[Sort[oneSets10[[k]]] ==
12      Sort[baselineOneSets10[[k]], {k, 1, 10}]

```

```

(* Connected input sets: *)
(* Node 1 (AND): {2,3} Node 6 (XNOR): {5,7} *)
(* Node 2 (OR): {1,3} Node 7 (NOT): {6} *)
(* Node 3 (XOR): {4,5} Node 8 (IMP): {1,9} *)
(* Node 4 (KOFN): {2,3,5} Node 9 (NIMP): {2,10} *)
(* Node 5 (NOR): {6} Node 10(MAJ): {3,4,7,8} *)
(* Per-node verification: *)
(* {True,True,True,True,True,True,True,True,True,True}*)
(* Out: True *)

```

All ten gates—AND, OR, XOR, KOFN, NOR, XNOR, NOT, IMPLIES, NIMPLIES, and MAJORITY—pass: the analytic one-sets coincide exactly with the exhaustive output columns produced by `ApplyGate`.

We then define six mixed queries—four full-output patterns and two subsystem projections—and compute their compressed representations via `mixedQueryRepresentation`, which intersects local one-sets and extracts the base set and offset family:

Listing 8: Mixed-query evaluation and corroboration.

```

1 queryIndices[nodes_, pattern_] := Sort@Fold[
2   Intersection, Range[2^10],
3   MapThread[
4     If[#2 == 1, oneSets10[[#1]],
5       Complement[Range[2^10], oneSets10[[#1]]]] &,
6     {nodes, pattern}]];
7
8 mixedQueryRepresentation[nodes_, pattern_] := Module[
9   {analytic, unionCoords, sumandos, baseIndices},
10  analytic= queryIndices[nodes, pattern];
11  unionCoords= Sort@DeleteDuplicates@
12    Flatten[ics10[[nodes]]];
13  sumandos= allOffsets[10, unionCoords];
14  baseIndices= Select[analytic,
15    Function[idx, AllTrue[
16      Complement[Range[10], unionCoords],
17      inputs10[[idx, #]] == 0&]]];
18  <|"DecimalRepertoire" -> baseIndices,
19    "Sumandos" -> sumandos|>;
20
21 resF1= mixedQueryRepresentation[Range[10],
22   {1,1,1,1,1,1,1,1,1,1}];
23 gpF1 = givePlaces[resF1["DecimalRepertoire"],
24   resF1["Sumandos"]];
25 Sort[gpF1] === Sort[queryIndices[Range[10],
26   {1,1,1,1,1,1,1,1,1,1}]]

```

```

(* F1 DecimalRepertoire: {143,215,399,400,471,472} *)
(* F1 Sumandos:          {0} *)
(* F1 Unfolded:          {143,215,399,400,471,472} *)
(* Out: True [verified for all six cases F1-F4,S1,S2]*)

```

Tables 3 and 4 show the results for all six cases. Every value shown is the actual output of the code above.

Case	Output pattern	DecimalRepertoire	Sumandos	Rows
F1	1111111111	{143, 215, 399, 400, 471, 472}	{0}	6
F2	1111111110	{15, 87, 271, 272, 343, 344}	{0}	6
F3	1111111101	{655, 727, 911, 912, 983, 984}	{0}	6
F4	1111101111	{79, 207, 335, 336, 463, 464}	{0}	6

Table 3: Compressed representation for the four full-output queries. Each case yields exactly 6 satisfying rows out of 1024. Because the full-output queries saturate all 10 coordinates ($C_q = [10]$, no free coordinates), the offset family is $\{0\}$ and the DecimalRepertoire is the complete one-set. This is the *coarse* form, in which one offset family is shared by every row; Section 4.2 gives the *fine* form, in which each row carries its own.

4.2 The merged description

Tables 3 and 4 give what we will call the *coarse* description: one base set, and a single offset family shared by every row of it, generated by the coordinates lying outside C_q . That description

Case	Queried nodes / pattern	DecimalRepertoire	Sumandos	Rows
S1	{4, 6, 7, 10} / 0111	{141, 217}	{0, 1, 256, 257, 512, 513, 768, 769}	16
S2	{4, 6, 7, 8, 9, 10} / 011101	{141, 217, 397, 398, 473, 474, 653, 729, 909, 910, 985, 986}	{0}	12

Table 4: Compressed representation for the subsystem queries. *S1* has free coordinates {1, 9, 10}, so the 2-element base set unfolds via 8 offsets to 16 rows. *S2* saturates all 10 coordinates, collapsing to a 12-element base set with trivial offsets. All indices verified exactly against the exhaustive baseline. As in Table 3, this is the coarse form; see Section 4.2.

is correct, and every index it produces was corroborated above. It is not, however, the shortest description of the same answer, because it frees a coordinate only when the coordinate is free for the *whole* query. A coordinate that is free for some rows but not others is an ordinary don’t-care, and the coarse form pays for it once per row.

The *fine* description gives each row its own offset family, which is what $\text{Dec}(L, \Omega)$ says in general: the sumandos are the fillings of a schema’s own don’t-care positions, wherever those positions fall, connected inputs included. Adjacent rows then merge. Applied to the six cases, the 38 base rows of Tables 3–4 become the 20 schemata of Table 5.

Listing 9: Merging the base set into schemata over C_q .

```

1 mergeSchemata[baseIndices_, unionCoords_, n_] :=
2   Module[{m = Length[unionCoords], minterms, dnf},
3     minterms = Table[
4       And @@ Table[
5         If[BitGet[j - 1, unionCoords[[i]] - 1] == 1,
6           cbv[i], Not[cbv[i]]],
7         {i, m}],
8       {j, baseIndices}];
9     dnf = BooleanMinimize[Or @@ minterms, "DNF"];
10    ... (* read each And term as fixed literals *);
11
12    (* each schema unfolds against ITS OWN offset family *)
13    unfoldSchema[s_, n_] := givePlaces[
14      {1 + Total[2^(s["activators"] - 1)]},
15      allOffsets[n, Join[s["activators"], s["inhibitors"]]]];

```

Two properties of this merge are worth stating precisely, because the obvious reading of “38 rows become 20” overstates both.

First, the 20 is a genuine minimum and not an artefact of the search. The schemata were computed twice, by implementations sharing no code: Wolfram’s `BooleanMinimize` and, independently, prime-implicant generation followed by a set cover. The two agree form for form in all six cases. Exhaustive search over the prime implicants confirms that no cover smaller than the one reported exists; in every case here each prime implicant is essential, so the cover is forced.

Second, and less comfortably, *a shorter cover is not the same thing as a shorter description*, and here the two objectives part company. A minimum cover minimises the number of schemata. Minimising bits is a different problem. Priced in a common coordinate—the same mask naming C_q , the same self-delimiting count, and then a payload of c_q symbols per object, at one bit for a base row and $\log_2 3 \approx 1.585$ bits for a schema coordinate—the merged form is shorter in only two of the six cases:

Case	Schemata over $\{1, \dots, 10\}$	Rows	Schemata
F1	01110001*0, 01101011*0, *111000110, *110101110	6	4
F2	01110000*0, 01101010*0, *111000010, *110101010	6	4
F3	01110001*1, 01101011*1, *111000111, *110101111	6	4
F4	0111001**0, *111001*10	6	2
S1	*0110001**, *0011011**	2	2
S2	00110001**, 00011011**, *01100011*, *00110111*	12	4
Total		38	20

Table 5: The fine description of the same six queries. Each string gives coordinates 1 to 10 left to right in the LSB-first convention used throughout, with * marking a don’t-care, so the string is a Holland schema and the * positions are precisely that schema’s own offset family. *Nothing about the network changes here.* Each case was checked to cover exactly the index set of the corresponding coarse row set, elementwise by symmetric difference, and against the exhaustive baseline. *S1* does not merge: its two rows already differ in more than one fixed coordinate.

Case	Coarse (bits)	Fine (bits)	Ratio	
F1–F3	75.0	78.4	1.05	longer
F4	75.0	44.7	0.60	shorter
S1	27.0	35.2	1.30	longer
S2	137.0	78.4	0.57	shorter
Total	464.0	393.5	0.85	

The trit payload is the cheaper of the two obvious schema codes—a code that transmits a fixed-position mask and then the values of the fixed positions costs $c_q + |\text{fixed}|$ bits, which is more for every schema above—so the comparison is already generous to the merged form. The pattern is readable: merging pays when it removes many rows and frees several coordinates (*S2*, $12 \rightarrow 4$; *F4*, $6 \rightarrow 2$), and costs when it removes few and frees one (*F1–F3*, $6 \rightarrow 4$ with a single * in each schema), or when nothing merges at all (*S1*). The aggregate saving of $0.85\times$ is real; the per-case claim that merging always shortens the description is not, and we do not make it.

4.3 Overlap analysis

The six query cases reveal the role of overlapping input supports. When several queried nodes share input coordinates, the effective search space shrinks from the sum of their local arities to the size of the coordinate union.

Define:

- $d_q = \sum_{i \in Q} |N_i|$: the sum of queried gate arities (the “raw” local-assignment space has size 2^{d_q}).
- $c_q = |C_q|$: the size of the union of their connected-input coordinates.
- $\mu_q = d_q - c_q$: the *overlap multiplicity*, counting repeated coordinate uses.
- $R_q = 2^{\mu_q}$: the exact reduction factor.

The most transparent case is *S1*. The four queried gates contribute $3 + 2 + 1 + 4 = 10$ local inputs, but their coordinate union is only $C_{S1} = \{2, 3, 4, 5, 6, 7, 8\}$ (seven coordinates). The

Case	Queried nodes	d_q	c_q	μ_q	R_q	C_q
F1–F4	$\{1, \dots, 10\}$	21	10	11	2048	$\{1, \dots, 10\}$
S1	$\{4, 6, 7, 10\}$	10	7	3	8	$\{2, 3, 4, 5, 6, 7, 8\}$
S2	$\{4, 6, 7, 8, 9, 10\}$	14	10	4	16	$\{1, \dots, 10\}$

Table 6: Overlap statistics for the mixed queries. The full-output queries have $d_q = 21$ (the sum of all local arities) but $c_q = 10$ (only 10 distinct coordinates exist), giving an overlap multiplicity of $\mu_q = 11$ and an exact reduction factor of $R_q = 2^{11} = 2048$.

three free coordinates $F_{S1} = \{1, 9, 10\}$ generate the offset family

$$\Omega(F_{S1}) = \{0, 1, 256, 257, 512, 513, 768, 769\},$$

and the zero-free base set is $L_{S1}^{(0)} = \{141, 217\}$. Deconvolution then gives

$$\mathcal{J}_{S1} = \text{Dec}(\{141, 217\}, \Omega(\{1, 9, 10\})),$$

which unfolds to the 16 corroborated rows. This is the exact multi-gate analogue of the single-gate AND deconvolution: overlap first shrinks the determining coordinate set from 10 to 7, and only then do the remaining free coordinates generate the full family of offsets.

4.4 What the compression reveals

Three structural observations emerge from the overlap analysis.

First, the full-output queries saturate all ten coordinates ($C_q = \{1, \dots, 10\}$), leaving no free coordinates. The offset family therefore collapses to the singleton $\{0\}$, and the DecimalRepertoire is the complete answer set. This is the limiting case of causal compression: the query is so tightly constrained that the base set alone suffices, and the analytic path reduces the raw 2^{21} -element local-assignment space to a direct enumeration of the few satisfying rows. The four cases F1–F4 each yield exactly six rows out of 1024, and their DecimalRepertoire sets differ precisely in those six positions—the topological skeleton is shared, but the gate-specific selection rules pick out different decimal anchors.

Second, even though the full-output queries have trivial offsets, they still benefit from overlap: the raw local-assignment space drops from 2^{21} to 2^{10} , an exact reduction factor of $2^{11} = 2048$, before any row materialisation is considered. The query does not pay once per gate for a shared coordinate; it pays once for the coordinate itself.

Third, what appears in the full repertoire as dispersed behaviour—six rows scattered across 1024—is explained analytically as overlap-aware reuse of shared coordinates. The compressed representation makes the generative structure visible, whereas the exhaustive table presents only the effect. The subsystem query S1 illustrates the complementary case: with three free coordinates, its two-element base set unfolds via eight offsets to 16 rows, showing how the offset family encodes the degrees of freedom that the query leaves unconstrained.

5 Scalability and Causal Compression

5.1 Resource-envelope benchmarks

The scalability of the analytic method depends on the support-union size $|C_q|$, not on the ambient network size n . To confirm this, a deterministic benchmark was executed across ring-local heterogeneous Boolean networks at $n \in \{30, 60, 80, 200\}$, with five independently seeded replicates per size. The gate catalogue covered all eleven non-canalising families. Three query tiers were measured:

- **T1** ($|Q| = 1$): a single queried node.
- **T2** ($|Q| = 4$): a four-node query.
- **T3** ($|Q| = 8$): an eight-node query (the primary stress case).

n	Query	$ C_q $	$n - C_q $	μ_q	Time (s)	Peak memory
30	T1	3	27	0	6.7×10^{-5}	4.3 KB
30	T2	7	23	4	2.4×10^{-4}	12.2 KB
30	T3	10	20	13	4.7×10^{-4}	14.4 KB
60	T1	3	57	0	6.4×10^{-5}	4.1 KB
60	T2	5	55	5	2.0×10^{-4}	8.7 KB
60	T3	10	50	10	9.7×10^{-4}	33.4 KB
80	T1	4	76	0	1.1×10^{-4}	7.6 KB
80	T2	6	74	4	1.9×10^{-4}	7.3 KB
80	T3	10	70	10	5.5×10^{-4}	18.2 KB
200	T1	3	197	0	6.2×10^{-5}	4.1 KB
200	T2	6	194	3	2.2×10^{-4}	8.8 KB
200	T3	10	190	10	1.1×10^{-3}	34.9 KB

Table 7: Executed exact causal evaluation measurements (median over five replicates per size). All measurements remain in the sub-millisecond to low-millisecond range across the full range $n = 30$ to $n = 200$, confirming that the cost is governed by $|C_q|$ rather than by n .

The central observation is that the median support-union size $|C_q|$ for the T3 tier remains approximately 10 across all four ambient sizes. As n grows from 30 to 200, the free-coordinate count grows from 20 to 190—the region the exact method avoids enumerating—while wall time fluctuates only within the sub-millisecond to low-millisecond range.

Table 8 contrasts the exact method against the naive exhaustive lower bounds.

n	Exact time (s)	Exact memory	Naive storage LB	Naive time LB
30	4.7×10^{-4}	14.4 KB	3.75 GB	≈ 1.07 s
60	9.7×10^{-4}	33.4 KB	7.50 EB	≈ 36.5 y
80	5.5×10^{-4}	18.2 KB	10.0 YB	$\approx 3.83 \times 10^7$ y
200	1.1×10^{-3}	34.9 KB	$\approx 4.02 \times 10^{61}$ B	$\approx 5.09 \times 10^{43}$ y

Table 8: Synthesis comparison for the T3 tier (median over five replicates). Naive storage lower bound is $n 2^n / 8$ bytes; naive time lower bound extrapolates at 10^9 rows per second. At $n = 60$ naive materialisation exceeds the exabyte threshold; at $n = 200$ it exceeds any physically realised medium by many orders of magnitude.

5.2 Description length and causal compression

A conceptually distinct comparison concerns the description length of the mechanism versus the complexity of the materialised output. The underlying principle, explored computationally in [3], is that algorithmically simple objects—those generated by short programmes—are *sensitive* to perturbation: modifying a single instruction in the generating programme produces non-linear changes in the output, because the programme’s parts are causally interdependent. By contrast, algorithmically random objects are *insensitive*: their shortest description is the object itself, and random perturbations change only a small fraction of the data. The ratio between the length of the generating programme and the length of its output is therefore a quantitative signature of causal structure.

The catalogue encoding. Because a description length is meaningless until the descriptive language is fixed, we state ours explicitly. We use two, and say at the outset which is primary. The first, D_{formula} , is a *length under the twelve-family catalogue*: it names a gate by its index in a fixed dictionary of twelve families. The second, D_{schema} , is catalogue-free: it writes the gate’s schemata out. D_{schema} is the measure we report as primary, for reasons given below and settled by measurement rather than preference. Each node i of a network of size n , carrying gate g_i with in-degree d_i , is encoded by three fields:

$$\underbrace{\log_2 K}_{\text{which gate}} + \underbrace{\log_2(n+1)}_{\text{in-degree } d_i} + \underbrace{\log_2 \binom{n}{d_i}}_{\text{which inputs}} + \underbrace{\pi(g_i, d_i, n)}_{\text{parameters}}, \quad K = 12,$$

and $D_{\text{formula}} = \sum_{i=1}^n$ of these costs. The parameter field π charges what each family actually requires: $\log_2(d_i+1)+1$ for KOFN (the threshold $k \in \{0, \dots, d_i\}$), $\log_2 n + 2$ for CANALISING (canalising index, canalising value, canalised output), $\log_2 \max(1, d_i(d_i-1))$ for IMPLIES and NIMPLIES (an *ordered* pair among the inputs), $\log_2 \max(1, d_i)$ for NOT, and a uniform one-bit polarity slot otherwise. Costs are idealised code lengths and are not rounded to whole bits.

The in-degree field is not decorative: without d_i a decoder can neither determine how many bits to read for the input set nor interpret them as an index into the d_i -subsets of $\{1, \dots, n\}$, so a code lacking it is not uniquely decodable and is therefore not a description length at all. The cost of specifying n itself is not charged, being constant across the networks compared here.

One property of this encoding must be stated plainly: it is a declared *choice*, not a canonical object. The invariance theorem guarantees only that a change of descriptive framework costs an additive constant, and in applications that constant can be arbitrarily large [14]. What the present setting secures is narrower and still worth stating: because the mechanism is given and the encoding is published, the constant is fixed and known rather than unknown and unbounded, so figures computed this way are exactly comparable across networks. D_{formula} is accordingly an upper bound on the shortest description in one declared language—never an estimate of K .

The catalogue-free encoding, and why it is primary. The first field above, $\log_2 K$ with $K = 12$, charges for a dictionary that is never transmitted. It is a fair price only for a decoder that already holds the catalogue, and it has a consequence that is easy to miss: under D_{formula} an XOR and an OR of the same in-degree cost *exactly the same*, because each is one entry in the dictionary. A measure that cannot separate those two is not responding to the object.

The schema normal form removes the dictionary. Each node is described by its schemata over the ambient coordinates—the same objects as Table 5, obtained the same way—and the code is a self-delimiting count of schemata followed, for each schema, by the number of fixed coordinates, which coordinates those are, and their values:

$$D_{\text{schema}} = \gamma(|S_i| + 1) + \sum_{s \in S_i} \left[\log_2(n+1) + \log_2 \binom{n}{k_s} + k_s \right],$$

summed over nodes, where S_i is node i ’s schema set and k_s the number of coordinates schema s fixes. Nothing is looked up. The code is decodable, and we verify that directly rather than asserting it: expanding the don’t-cares of every schema regenerates each node’s one-set exactly, elementwise against exhaustive evaluation, for all ten nodes of the benchmark.

D_{schema} is the larger number, and that is the point. It is the price of not charging for a dictionary one did not send. It also separates gates that D_{formula} cannot: the OR is cheap because its minterms merge into few templates, and the XOR expensive because they merge into none.

Measured values. For the 10-node benchmark:

Measure	Value	Interpretation
Formula components C_{formula}	23	symbolic pieces in the analytic description
Programme length D_{schema}	232.72 bits	the mechanism, catalogue-free (primary)
Programme length D_{formula}	135.66 bits	the same mechanism, under the twelve-family catalogue
BDM of the output	580.01 bits	algorithmic estimate of the behaviour
ZIP-compressed output	10 016 bits	lossless compression of the output table
Shannon content H_{total}	10 229.61 bits	repertoire-scale i.i.d. coding
BDM/ D_{schema}	2.49	estimator within a small factor of the truth
ZIP/ D_{schema}	43.04	mechanism smaller by a factor of ≈ 43
$H_{\text{total}}/D_{\text{schema}}$	43.96	mechanism is far shorter than full coding

The exact formulae are not merely another way to print the same table: they constitute a dramatically shorter generative description. D_{schema} measures the cost of specifying the mechanism; ZIP and H_{total} measure the cost of describing what that mechanism produces. The load-bearing comparison is with ZIP, since a general-purpose compressor is a genuine adversary: it is free to exploit any regularity it can find in the realised outputs, and it still requires a factor of ≈ 43 more bits than the generative law.

We state that factor as a factor and not as “two orders of magnitude”, because the catalogue-free measure does not support the rounder phrase: under D_{formula} the ratio to ZIP is 73.8, or 1.87 orders of magnitude, and under D_{schema} it is 43.04, or 1.63. The qualitative claim—that the generative description is smaller than any statistical encoding of the behaviour by well over an order of magnitude—survives the change of language. The wording “two orders of magnitude” does not, and we have removed it throughout.

H_{total} deserves a caveat rather than emphasis. At 10 229.61 bits it is 99.9 per cent of the raw 10 240-bit table, because the output bits are close to balanced and an entropy code therefore saves almost nothing. Far from weakening the result, this is the sharpest way to state it: the behaviour of this network is *statistically near-featureless* and nevertheless generated by a description of a hundred bits. Statistical featurelessness and extreme algorithmic compressibility coexist here, which is precisely the regime in which distributional summaries are least informative and generative descriptions most so.

Calibrating an algorithmic estimator against known ground truth. Comparing a generative law only against statistical baselines would be a weak test, since it selects opponents that are constitutionally unable to see generative structure. The stronger comparison is with a method designed to estimate algorithmic content without knowing the mechanism. We therefore compute the Block Decomposition Method (BDM) [14] over the same $2^{10} \times 10$ output repertoire. Because the generating programme and its exact length are both known here, this is a rare setting: the estimator can be assessed against ground truth rather than against another estimate.

BDM returns 580.01 bits, against 10 016 for ZIP and 10 229.61 for H_{total} . The algorithmic estimator is thus roughly seventeen times more efficient than lossless compression on an object where compression is nearly powerless, which is a substantial vindication of the algorithmic approach over the statistical one. Two controls confirm that BDM is responding to the generative structure and not to the marginals: a row permutation, which preserves every row of the table and destroys only the ordering geometry, raises BDM to 14 714 bits; and a density-matched

random matrix gives 15 545 bits. The separations are 25.4-fold and 26.8-fold respectively, and they are robust to the choice of partition strategy.¹

The residual gap is the quantity of interest. At 580.01 bits against a true programme length of 232.72, BDM sits within a factor of 2.49 of the exact description—a figure that means little on its own, so we give it its reference distribution in the same breath: across 200 random networks at the same $n = 10$, $\text{BDM}/D_{\text{schema}}$ has median 2.34 and range 0.86 to 5.03, so the benchmark sits close to the middle of its own class rather than being a favourable instance. That factor is a measure of the price of not knowing the generator: it is what an estimator must pay when the mechanism has to be recovered from the output rather than read off from the specification. It is small, which speaks well of the estimator, and it is not zero, which is what one should expect of any method that must infer rather than compute.

Do the two sides actually track each other? A single ratio is not a relationship, and the natural way for one to be spurious is that both quantities grow with the size of the object. We remove that explanation by construction rather than adjusting for it afterwards: the 200 networks all have $n = 10$, so every output object is a 1024×10 matrix, with identical shape, cell count and BDM partition. Size does not vary, so it cannot be the reason for anything below. Only the mechanism—gate families, in-degrees and wiring—varies, drawn at random from the twelve families.

Across that sample D_{schema} and BDM are positively but only moderately associated: Pearson $r = +0.388$ and Spearman $\rho = +0.424$, against a permutation null on the pairing that holds both marginals exactly and has central 95 per cent interval $[-0.136, +0.144]$ over 10^4 shuffles ($p = 10^{-4}$). The association is real and it is not proportionality; D_{schema} accounts for roughly fifteen per cent of the variance in BDM. We claim no more than that, and in particular the factor above should be read as one measurement within a spread, not as a constant.

Two controls decide whether the catalogue-free measure is doing any work. The first is D_{formula} , which never reads an output bit; the second is the raw degree budget $\sum_i d_i$, the crudest mechanism summary available. Against BDM these give $r = +0.207$ and $r = +0.245$ respectively—so the catalogue measure is, on this sample, a slightly *worse* predictor of the behaviour than simply counting edges. Regressing the degree budget out of both sides separates them cleanly: D_{schema} retains $r = +0.311$ (null 95 per cent $[-0.135, +0.139]$, $p = 2 \times 10^{-4}$) while D_{formula} falls to $r = +0.020$, which is to say nothing at all. The whole of D_{formula} ’s association with the behaviour is the wiring budget; D_{schema} has signal beyond it. That is the quantitative form of the point made above—the catalogue charges $\log_2 12$ for XOR and OR alike, and so cannot see the difference that BDM sees—and it is the reason D_{schema} is reported as primary. The pattern is stable under our own knobs: across seeds and in-degree ceilings of 3, 4 and 5, $r(D_{\text{schema}}, \text{BDM})$ stays in $[+0.317, +0.430]$ and $r(D_{\text{formula}}, \text{BDM})$ in $[+0.114, +0.207]$.

Finally, the direction is as theory requires but not universally. For a deterministic system $K(\text{output}) \leq K(\text{mechanism}) + O(1)$, so BDM should exceed D_{schema} , the overshoot being block decomposition’s known tendency to charge cross-block structure once per occurrence rather than a defect of either quantity. It does so in 194 of the 200 networks. The six exceptions are not suppressed: D_{schema} is an upper bound on the mechanism’s description length, not the mechanism’s Kolmogorov complexity, and BDM is an estimate rather than a bound, so nothing forbids the ordering from reversing on a network whose schema normal form is verbose relative to what the estimator recovers.

Where the remaining gap comes from is also visible. Block decomposition combines local estimates with a global account of repetition, so structure that spans blocks must be paid for once

¹All values use a recursive partition covering every column. The default block partition in common implementations discards leftover columns—at $n = 10$ it silently measures eight of the ten—and we report the recursive figure throughout. BDM is also known to converge toward Shannon entropy once block sizes outrun the available CTM tables, so these values should be read as scale-dependent estimates.

per occurrence; recent work extends the method with reusable program code and conditional descriptions precisely to charge shared structure only once [8]. The representation studied here is the limiting case of that idea. The base set L is the code, written once; the offset family Ω is the schedule of its reuse, and because Ω factorises as a sumset of the free weights the schedule is itself given in closed form rather than discovered. An estimator must infer both the code and the schedule from the output; here the mechanism supplies them. That is the whole of the difference between 232.72 and 580.01 bits, and it is why the comparison is a calibration rather than a competition. Appendix B applies the same estimator at the level of a single gate definition, where it serves as a check on the representation rather than on the method.

This is causal compression in a precise, non-rhetorical sense: the 23 formula components form a short programme whose output spans 1024 rows and 206 distinct output states.

A single experiment fixes what the two sides of that statement do and do not measure, and forestalls the most natural misreading. Rewire node 10 to a different input set of the same cardinality. The behaviour changes substantially: 440 of the 1024 output rows differ, and the number of distinct output states falls from 206 to 172. Neither description length changes at all. D_{formula} is bit-identical, because it is a function of n , the in-degrees and the gate types alone and never reads an output bit; and so, we should be equally plain, is D_{schema} , whose fields depend only on n and on the schema set that the gate and its in-degree determine. Over 200 random rewirings that preserve every gate and every in-degree, D_{schema} takes exactly *one* distinct value, 232.72 bits, while the behaviour moves in a median of 1012 of the 1024 rows.

Both facts are consequences of the same design, and promoting D_{schema} does not soften the caveat. Each is the correct quantity for the programme side of a programme-versus-output comparison, and neither is a complexity measure of behaviour: they cannot rank networks by the richness of what those networks do, and no claim of that form is made here. This is precisely why the correlation reported above is with the *gate composition* rather than with the wiring—rewiring is the axis along which D_{schema} is deliberately blind. The behavioural side of the comparison is carried by BDM, ZIP and H_{total} , which respond to the realised output and not to the specification.

5.3 Validation summary

Exact reconstruction has been verified at $n = 10$ and $n = 13$ by exhaustive comparison against the baseline repertoire, and beyond the exhaustive-baseline range at $n \in \{16, 20\}$ by a stratified closed-form-set audit across all twelve gate families (seeds 301 and 302): each node’s packaged closed-form one-set is materialised and compared entry-by-entry against the full gate truth table over all 2^d connected-input patterns, and the composed analytic prediction is checked row-wise against direct evaluation on 1020 Hamming-stratified rows per seed, with zero elementwise mismatches in all cases. A dispatch-consistency audit additionally covers $n \in \{20, 50\}$ (1020 stratified rows, seeds 201 and 202); it validates large-scale evaluation plumbing and carries no theorem weight.

6 Dynamical Landscape

The previous sections answered an exact inverse question: which input rows produce selected mixed output events. We now add the complementary forward-dynamical question on the same 10-node network: which output patterns are themselves one-step realisable, and which belong to the recurrent part of the state-transition graph?

Let $F : \{0, 1\}^{10} \rightarrow \{0, 1\}^{10}$ be the synchronous update map. Exhaustive enumeration of all $2^{10} = 1024$ states gives

$$|\text{Im}(F)| = 206, \quad |\mathcal{U}_{10} \setminus \text{Im}(F)| = 818, \quad \frac{|\text{Im}(F)|}{|\mathcal{U}_{10}|} \approx 0.201.$$

Claim	Validation layer	Evidence	Scale
Gate one-set exactness	Gate algebra; property tests	Per-gate corroboration	$ N_i \leq 6$
Ordering invariance	Six-node check	Involution transport	$n = 6$
Network exactness	Dispatch consistency	Zero discrepancy; accuracy 1.0	$n \leq 13$
Cost governed by $ C_q $	Resource-envelope benchmark	Tables 7–8	$n \leq 200$
Sampling reduces validation cost	Stratified sampled audit	Accuracy 1.0; seeds 301, 302	$n \in \{20, 50\}$

Table 9: Mapping from claims to validation evidence.

Only about 20.1% of all 10-bit patterns are admissible as one-step outputs. The local logical constraints and the overlap structure restrict the forward image to a small structured subset of the full configuration space.

The recurrent dynamics collapse further. Table 10 shows that the entire functional graph has only four attractors: one fixed point, two 2-cycles, and one 6-cycle.

Attractor	Period	Basin size	Recurrent states
A_1	1	488	$\{0000010100\}$
A_2	2	320	$\{1101010110, 0110010110\}$
A_3	6	204	$\{1111010101, 1111010001, 1111010000, 1111010010, 1111010110, 1111010111\}$
A_4	2	12	$\{1111010100, 1111010011\}$

Table 10: Exact recurrent structure of the 10-node benchmark. Basin sizes sum to 1024; the three largest basins cover 98.8% of the state space.

The connection between the exact causal queries and the dynamical landscape is immediate. The four full-output patterns $F1$ – $F4$ are all one-step reachable states, but none is itself an attractor state. Cases $F1$ – $F3$ converge to the 2-cycle A_2 within three steps; $F4$ converges to the fixed point A_1 within six steps. The subsystem query $S1$ induces five distinct image states spread across all four basins; adding the extra $S2$ constraints narrows this to three image states in three basins, removing the branch to A_4 . Stronger exact causal constraints reduce not only combinatorial freedom but also dynamical dispersion.

The causal descriptions (backward query) and the dynamical landscape (forward iteration) are complementary views of the same network. The backward query identifies the generative structure—which input conditions produce a given event. The forward dynamics reveal the eventual fate of those events under iterated application of the same causal rules. Neither view alone exhausts the network’s behaviour; together they provide both mechanism and consequence.

In the language of AID, the backward query is a perturbation test: the system is interrogated about its own computational capabilities, and the compressed representation is its answer. The forward dynamics show that this answer has further consequences: the six rows satisfying $F1$ do not scatter randomly through the state-transition graph but converge to a specific attractor within three steps. The causal rules that determine *which* input conditions produce a given output also determine *where* that output leads under iteration—the mechanism governs both the generative and the consequential structure of the network.

7 Discussion

The central contribution of this paper is the formalisation and computational validation of a causal compression methodology for Boolean networks. The methodology rests on the observation that deterministic local rules imprint self-similar patterns across the exhaustive repertoire, and that these patterns admit exact short descriptions. We discuss the principal results and their connections to the broader programme of algorithmic causality.

General formulae for all gate families. The self-similar patterns first observed empirically for AND, OR, and XOR gates [3] turn out to have a uniform algebraic structure that holds across all twelve canonical families. Each gate type determines a characteristic selection rule over the connected-input assignments—band intersections for monotone gates, parity classes for XOR/XNOR, Hamming-weight layers for threshold gates, asymmetric pairs for implication, hierarchical collapse for canalising gates—while the part of the offset family contributed by the disconnected coordinates is identical for all gates sharing the same topology. Only that part is purely topological. The remainder of Ω is gate-dependent, arising from connected inputs that a given schema leaves free, and it is exactly what distinguishes families whose one-sets merge into few schemata from those whose do not. With that qualification the separation is not a coincidence: it reflects the factorisation of the network’s causal structure into a gate-semantic component (the base set) and a topological component (the disconnected part of the offset family). Together they constitute a short generative programme for the full one-set of any gate.

Overlap compression. When multiple nodes are queried simultaneously, shared input coordinates reduce the effective search space by an exact factor of 2^{μ_q} . For the 10-node benchmark, this factor reaches 2048 for the full-output queries. The compressed representation makes this reduction explicit: the base set encodes the constrained coordinates, the offset family encodes the free coordinates, and deconvolution recovers the full answer. The full-output queries illustrate the limiting case: all coordinates are constrained, the offset family collapses to $\{0\}$, and the DecimalRepertoire alone constitutes the complete answer.

Scalability. The analytic evaluation cost is governed by $|C_q|$, not by n . Under bounded in-degree and bounded query scope, median analytic evaluation stays at or below roughly a millisecond at $n = 200$, a regime where naive exhaustive materialisation is physically impossible; the salient fact is not the absolute figure, which is hardware-dependent, but that it is indistinguishable from the figure at $n = 30$. In the language of Section 2.2 this is a statement about schema order: what must be evaluated is the number of fixed positions, not the size of the space in which they sit.

The method as the exactly solvable limit of AID. The methodology realises the AID principle [11, 13, 12] that short generative descriptions of complex behaviour can be identified by exploiting the causal structure of the generating mechanism—but it does so in the restricted regime where the mechanism is given rather than inferred. The compressed representation is the “short programme”; the full repertoire is the “complex output”; and the factor of ≈ 43 between D_{schema} and the cost of describing the behaviour is a quantitative measure of causal compression.

Because the mechanism is known here, no estimator is needed and none is proposed. The methodological contribution runs in the other direction: the networks studied here are objects whose exact generating programme and exact description length are both available, which makes them a calibration class for methods that must operate without that knowledge. Section 5.2 uses them so, and the outcome is informative in a way that a purely internal comparison could not be—the algorithmic estimator recovers the structure that statistical measures miss entirely, and lands within a small constant factor of the true programme length.

The connection to perturbation analysis is direct: a system whose causal description is short is one whose output changes non-linearly under perturbation of the generating programme [3], as the rewiring experiment of Section 5.2 shows concretely. This sensitivity is the hallmark of an algorithmically simple, causally integrated system—precisely the regime in which the

methodology is most powerful. The density-matched random control of the same section is the opposite pole: an object with no short causal description, on which the method would offer no compression advantage and on which, as measured, the algorithmic estimator itself returns a value close to the raw size.

Connection to Wolfram. The Principle of Computational Equivalence [9] holds that simple rules generate rich behaviour. The present method makes this quantitative: the 10-node network is specified by 23 formula components, yet its full behaviour spans 1024 rows with 206 distinct output states, 4 attractors, and structured basin geometry. The causal description is short; the generated behaviour is complex.

Limitations. The computational advantage depends on bounded support size: when $|C_q|$ approaches n , the analytic path converges to the cost of exhaustive materialisation. The method is therefore most powerful for networks with bounded in-degree and localised query scope. The present formalism treats synchronous one-step update; asynchronous or stochastic schedules are a direction for future work.

Three qualifications attach to the description-length results specifically. First, both D_{schema} and D_{formula} are lengths under declared encodings, upper bounds on the shortest description in their respective languages and not estimates of K ; a different but equally reasonable encoding would shift either by a constant. That the two languages give 232.72 and 135.66 bits for the same network is the plainest available illustration of the point, and is why we report both rather than only the smaller. Second, the calibration of Section 5.2 is anchored on a single network, and while we now place it within a 200-network reference distribution at the same size, that distribution is confined to $n = 10$; the factor of 2.49 should be read as one measurement within a spread of 0.86 to 5.03 rather than as a general constant, and how it behaves across sizes and in-degree regimes remains open. Establishing that across network families is the natural next study, and the present method supplies the ground truth such a study would require. Third, the estimator itself is scale-dependent, converging toward Shannon entropy once block sizes outrun the available base tables, so the comparison is meaningful in the regime reported and should not be extrapolated beyond it without re-measurement.

8 Conclusion

We have presented a causal compression methodology for Boolean networks that derives exact short descriptions of network behaviour from the interaction between gate semantics and network topology, validated computationally across all twelve canonical gate families, overlap-aware multi-node queries, and networks of up to $n = 200$ nodes.

The method rests on a uniform decomposition: every gate family determines a base set over its connected inputs, and every topology generates an offset family from its disconnected coordinates. Deconvolution of this pair recovers the exact one-set without row-wise gate evaluation. When multiple nodes are queried, overlapping input supports compress the effective search space by an exact factor of 2^{μ_q} .

The computational characterisation confirms that analytic evaluation remains of the order of a millisecond at $n = 200$ under bounded in-degree, and does not grow with n , that the causal description ($D_{\text{schema}} = 232.72$ bits) is shorter than any statistical encoding of the realised outputs by a factor of ≈ 43 (ZIP = 10 016 bits, $H_{\text{total}} = 10\,229.61$ bits) while a block-decomposition estimator, given no access to the mechanism, recovers the structure to within a factor of 2.49, and that all results are verified exactly against exhaustive baselines at tractable sizes.

The causal descriptions (backward query) and the dynamical landscape (forward iteration) provide complementary views of the same network. Together, they demonstrate that complex Boolean-network behaviour can be represented and reconstructed by short exact causal formulae. The compressed representation is not merely a computational shortcut; it identifies the generative structure of the mechanism—the short programme whose execution produces the ob-

served behaviour. This is causal compression in the precise sense of the Algorithmic Information Dynamics programme [11]: the network admits a description that is dramatically shorter than its output, and this brevity is a direct consequence of the deterministic causal rules that generate it.

The method provides a general tool for any domain in which Boolean networks model discrete dynamical systems. Future directions include extension to asynchronous update schedules, application to biological network models, and integration with the perturbation-based causal discovery pipeline of AID.

All computational experiments are reproducible via the companion code repository.

A Companion Code and Reproducibility

All computational results reported in this paper are reproducible using a companion code repository, available at the DOI listed on the title page. The repository is organised into self-contained modules corresponding to the main experimental sections of the paper:

1. **Six-node corroboration** (Section 3.1): AND gate deconvolution on the six-node network, with exhaustive verification against the gate-dispatch baseline.
2. **Ten-node mixed-gate benchmark** (Section 4.1): definition of the heterogeneous 10-node network, per-node analytic one-set computation for all twelve gate families, and the six mixed-query cases (F1–F4, S1, S2) with overlap analysis.
3. **Dynamical landscape** (Section 6): exhaustive enumeration of the state-transition graph, attractor identification, and basin-size computation for the 10-node benchmark.
4. **Resource-envelope benchmarks** (Section 5.1): scalability measurements at $n \in \{30, 60, 80, 200\}$ with five independently seeded replicates per size, covering three query tiers (T1, T2, T3).
5. **Description-length analysis** (Section 5.2): computation of D_{schema} , D_{formula} , C_{formula} , ZIP-compressed output length, and H_{total} for the 10-node benchmark, together with the 200-network BDM comparison and its permutation null.

Each module is self-contained and produces deterministic output. The analytic computations are implemented in Wolfram Language (requiring WolframKernel 13.0 or later); the scalability and description-length analyses are implemented in Python 3.9+ using only the standard library. The BDM calibration of Section 5.2 additionally requires `numpy` and a BDM implementation supplying the CTM base tables; the partition strategy and all seeds are recorded in the script. A shared library provides the core functions `ApplyGate` and `CreateRepertoiresDispatch` used throughout.

A replication notebook, `replication_comp_paper.ipynb`, re-executes every Python module and checks each numerical claim in this paper against the value printed here, reporting a consolidated pass/fail table and failing loudly if any result ceases to reproduce. It is accompanied by `D_formula_explained.ipynb`, a didactic companion that derives D_{formula} from first principles. The notebooks require `numpy`; all seeds are pinned and recorded.

B Algorithmic Complexity of the Gate Definitions

Section 5.2 measures the description length of a whole network and the algorithmic content of the behaviour it generates. A separate and narrower question is how much algorithmic content resides in the *definition* of an individual gate. We record the measurement here because it bears on a methodological choice that is easy to get wrong.

Representation. A gate must be presented to an algorithmic measure in some binary form, and the choice is not innocent. The form we use is the gate’s *truth table* at a fixed arity: for in-degree d this is a 2^d -bit string, it is the gate’s extensional definition, it is canonical, and it has the same length for every family, so values are directly comparable. Table 11 reports BDM over these strings at $d = 4$.

Gate	Truth table ($d = 4$)	Ones	BDM
XOR	0110100110010110	8	41.540
XNOR	1001011001101001	8	41.540
KOFN ($k=2$)	0001011101111111	11	41.137
CANALISING	0111111100000000	7	40.150
NOT	1111111100000000	8	38.220
IMPLIES	1111111100001111	12	38.220
NIMPLIES	0000000011110000	4	38.220
MAJORITY	0000000100010111	5	38.166
OR	0111111111111111	15	34.977
NOR	1000000000000000	1	34.977
AND	0000000000000001	1	33.797
NAND	1111111111111110	15	33.797

Table 11: BDM over truth tables at arity four, in bits, ordered by value. Parity gates rank above threshold gates, which rank above monotone gates. Every complement pair receives an identical value.

Two properties of the table are checks rather than results, and both pass. Parity gates rank above threshold gates, which in turn rank above the monotone gates—the expected ordering, since an AND truth table is a single one among fifteen zeros while an XOR table alternates. And each complement pair (AND/NAND, OR/NOR, XOR/XNOR) receives an *identical* value, which is correct: complementation is an $O(1)$ operation, so it cannot change algorithmic content. A measure that separated such pairs, or that ranked AND above XOR, would be responding to something other than logical structure.

Two representations that must not be used. The same measurement was attempted under two other binarisations and both fail, for reasons worth recording because the failures are instructive rather than incidental.

Assigning each family an arbitrary fixed-width code and measuring the resulting stream reports a property of the *labelling*, not of the gates: three arbitrary code assignments over the identical ten-node network give 104.58, 101.91 and 106.82 bits. The network did not change.

Measuring the ASCII bits of the gate names is worse. BDM then tracks the length of the English word almost exactly—the correlation between value and name length across the twelve families is $r = 0.998$ —so the measure would rank MAJORITY as the most complex gate because its name is the longest.

Both failures are instances of the same fault, and it is the fault that Section 5.2 warns against in the case of D_{formula} , and which is the reason D_{schema} is reported as primary: a measure must respond to the object and not to the representation chosen for it. Charging $\log_2 12$ for a dictionary entry is a milder version of charging for the length of a name. The truth table is used here precisely because it is the one presentation of a gate that is fixed by the gate itself.

References

- [1] Edmund M. Clarke et al. A compositional framework for boolean networks. *BioSystems*, 2019. The web result summarizes a formal compositional framework for Boolean networks.

- [2] Alberto Hernández-Espinosa. Causal index-set calculus for Boolean networks and deterministic reconstruction. Manuscript in preparation, 2026.
- [3] Alberto Hernández-Espinosa, Hector Zenil, Narsis A. Kiani, and Jesper Tegnér. Estimations of integrated information based on algorithmic complexity and dynamic querying. *Entropy*, 21(10):947, 2019. arXiv:1904.10393.
- [4] John H. Holland. *Adaptation in Natural and Artificial Systems: An Introductory Analysis with Applications to Biology, Control, and Artificial Intelligence*. University of Michigan Press, Ann Arbor, MI, 1975.
- [5] Stuart A. Kauffman. Metabolic stability and epigenesis in randomly constructed genetic nets. *Journal of Theoretical Biology*, 22(3):437–467, 1969.
- [6] Stuart A. Kauffman. *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford University Press, New York, 1993.
- [7] Ting Li et al. Dynamics of boolean networks: An exact solution. *Physical Review Letters*, 106(21):214101, 2011.
- [8] Eduardo Yuji Sakabe, Felipe S. Abrahão, Santiago Hernández-Orozco, Ricardo Gudwin, and Hector Zenil. Tighter bounds for algorithmic complexity estimation using a reusable code-based block decomposition method, 2026.
- [9] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, Champaign, IL, 2002.
- [10] Hector Zenil. Compression is comprehension: The unreasonable effectiveness of digital computation in the natural world. In *Unravelling Complexity: The Life and Work of Gregory Chaitin*. World Scientific, 2020. Preprint available as arXiv:1904.10258.
- [11] Hector Zenil, Narsis A. Kiani, Francesco Marabita, Yue Deng, Szabolcs Elias, Angelika Schmidt, Gordon Ball, and Jesper Tegnér. An algorithmic information calculus for causal discovery and reprogramming systems. *iScience*, 19:1160–1172, 2019.
- [12] Hector Zenil, Narsis A. Kiani, and Jesper Tegnér. *Algorithmic Information Dynamics: A Computational Approach to Causality with Applications to Living Systems*. Cambridge University Press, Cambridge, 2023.
- [13] Hector Zenil, Narsis A. Kiani, Allan A. Zea, and Jesper Tegnér. Causal deconvolution by algorithmic generative models. *Nature Machine Intelligence*, 1(1):58–66, 2019.
- [14] Hector Zenil, Fernando Soler-Toscano, Narsis A. Kiani, Santiago Hernández-Orozco, and Antonio Rueda-Toicen. A decomposition method for global evaluation of Shannon entropy and local estimations of algorithmic complexity. *Entropy*, 20(8):605, 2018.